

Experiment 7 :- To control network traffic using iptables rules.

Name: Vedant Sandip Kannurkar

UID: 2024801005

Division: B2

Exercise 1: iptables Essentials:

Steps performed :

View all rules in the filter table : iptables -L

```
admin4@gordon-freeman:~$ sudo iptables -L
[sudo] password for admin4:
Chain INPUT (policy ACCEPT)
target     prot opt source          destination
DOCKER-USER  all  --  anywhere       anywhere
DOCKER-ISOLATION-STAGE-1  all  --  anywhere       anywhere
ACCEPT     all  --  anywhere       anywhere           ctstate RELATED,ESTABLISHED
ACCEPT     all  --  anywhere       anywhere
ACCEPT     all  --  anywhere       anywhere
ACCEPT     all  --  anywhere       anywhere

Chain FORWARD (policy DROP)
target     prot opt source          destination
ACCEPT     all  --  anywhere       anywhere           ctstate RELATED,ESTABLISHED
ACCEPT     all  --  anywhere       anywhere
ACCEPT     all  --  anywhere       anywhere
ACCEPT     all  --  anywhere       anywhere

Chain OUTPUT (policy ACCEPT)
target     prot opt source          destination

Chain DOCKER (1 references)
target     prot opt source          destination

Chain DOCKER-ISOLATION-STAGE-1 (1 references)
target     prot opt source          destination
DOCKER-ISOLATION-STAGE-2  all  --  anywhere       anywhere
RETURN    all  --  anywhere       anywhere

Chain DOCKER-ISOLATION-STAGE-2 (1 references)
target     prot opt source          destination
DROP      all  --  anywhere       anywhere
RETURN    all  --  anywhere       anywhere

Chain DOCKER-USER (1 references)
target     prot opt source          destination
RETURN    all  --  anywhere       anywhere
```

Verbose view : iptables -L -v

Experiment 7 :- To control network traffic using iptables rules.

```
admin4@gordon-freeman:~$ sudo iptables -L -v
Chain INPUT (policy ACCEPT 0 packets, 0 bytes)
 pkts bytes target     prot opt in     out     source          destination
Chain FORWARD (policy DROP 0 packets, 0 bytes)
 pkts bytes target     prot opt in     out     source          destination
  0      0 DOCKER-USER  all   --  any    any    anywhere       anywhere
  0      0 DOCKER-ISOLATION-STAGE-1 all   --  any    any    anywhere       anywhere
  0      0 ACCEPT     all   --  any    docker0 anywhere       anywhere
  0      0 DOCKER     all   --  any    docker0 anywhere       anywhere
  0      0 ACCEPT     all   --  docker0 !docker0 anywhere       anywhere
  0      0 ACCEPT     all   --  docker0 docker0  anywhere       anywhere

Chain OUTPUT (policy ACCEPT 0 packets, 0 bytes)
 pkts bytes target     prot opt in     out     source          destination
Chain DOCKER (1 references)
 pkts bytes target     prot opt in     out     source          destination
Chain DOCKER-ISOLATION-STAGE-1 (1 references)
 pkts bytes target     prot opt in     out     source          destination
  0      0 DOCKER-ISOLATION-STAGE-2 all   --  docker0 !docker0 anywhere       anywhere
  0      0 RETURN     all   --  any    any    anywhere       anywhere

Chain DOCKER-ISOLATION-STAGE-2 (1 references)
 pkts bytes target     prot opt in     out     source          destination
  0      0 DROP       all   --  any    docker0 anywhere       anywhere
  0      0 RETURN     all   --  any    any    anywhere       anywhere

Chain DOCKER-USER (1 references)
 pkts bytes target     prot opt in     out     source          destination
  0      0 RETURN     all   --  any    any    anywhere       anywhere
```

To display only the rules under the INPUT chain:

```
admin4@gordon-freeman:~$ sudo iptables -v -L INPUT
Chain INPUT (policy ACCEPT 0 packets, 0 bytes)
 pkts bytes target     prot opt in     out     source          destination
```

To display all the rules under the mangle table:

```
admin4@gordon-freeman:~$ sudo iptables -L -t mangle
Chain PREROUTING (policy ACCEPT)
target     prot opt source          destination
Chain INPUT (policy ACCEPT)
target     prot opt source          destination
Chain FORWARD (policy ACCEPT)
target     prot opt source          destination
Chain OUTPUT (policy ACCEPT)
target     prot opt source          destination
Chain POSTROUTING (policy ACCEPT)
target     prot opt source          destination
```

To display all the rules under the nat table:

Experiment 7 :- To control network traffic using iptables rules.

```
admin4@gordon-freeman:~$ sudo iptables -L -t nat
Chain PREROUTING (policy ACCEPT)
target    prot opt source          destination
DOCKER    all  --  anywhere       anywhere            ADDRTYPE match dst-type LOCAL

Chain INPUT (policy ACCEPT)
target    prot opt source          destination

Chain OUTPUT (policy ACCEPT)
target    prot opt source          destination
DOCKER    all  --  anywhere       !localhost/8        ADDRTYPE match dst-type LOCAL

Chain POSTROUTING (policy ACCEPT)
target    prot opt source          destination
MASQUERADE all  --  172.17.0.0/16 anywhere

Chain DOCKER (2 references)
target    prot opt source          destination
RETURN   all  --  anywhere       anywhere
```

To flush all the current rules:

```
admin4@gordon-freeman:~$ sudo iptables --flush
admin4@gordon-freeman:~$ sudo iptables -N mychain
```

To create your own chain:

```
admin4@gordon-freeman:~$ sudo iptables -N mychain
admin4@gordon-freeman:~$ sudo iptables -L mychain
Chain mychain (0 references)
target    prot opt source          destination
```

To delete chains

```
admin4@gordon-freeman:~$ sudo iptables -X INPUT
iptables: No chain/target/match by that name.
admin4@gordon-freeman:~$ sudo iptables -X mychain
admin4@gordon-freeman:~$ sudo iptables -L mychain
iptables v1.8.7 (nf_tables): chain `mychain' in table `filter' is incompatible, use 'nft' tool.
```

Exercise 2: Basic Packet Filtering

To filter ICMP packets types: Before starting ensure that you can ping your partner system and that your partner system can ping you too successfully.

Experiment 7 :- To control network traffic using iptables rules.

```
admin4@gordon-freeman:~$ sudo ping -c 2 serverPR
ping: serverPR: Temporary failure in name resolution
admin4@gordon-freeman:~$ sudo ping -c 2 www.google.com
PING www.google.com (142.251.220.4) 56(84) bytes of data.
64 bytes from pnbomb-ay-in-f4.1e100.net (142.251.220.4): icmp_seq=1 ttl=117 time=1.63 ms
64 bytes from pnbomb-ay-in-f4.1e100.net (142.251.220.4): icmp_seq=2 ttl=117 time=10.2 ms

--- www.google.com ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 5074ms
rtt min/avg/max/mdev = 1.633/5.913/10.194/4.280 ms
```

Flush all your existing rules

```
admin4@gordon-freeman:~$ sudo iptables -F
```

Create a rule to prevent all outgoing icmp type packets to any destination. Test it using ping

```
admin4@gordon-freeman:~$ sudo iptables -A OUTPUT -o enp2s0 -p icmp -j DROP
admin4@gordon-freeman:~$ sudo ping -c 2 serverPR
ping: serverPR: Temporary failure in name resolution
admin4@gordon-freeman:~$ sudo ping -c 2 www.google.com
PING www.google.com (142.251.220.4) 56(84) bytes of data.
^C
--- www.google.com ping statistics ---
2 packets transmitted, 0 received, 100% packet loss, time 1056ms
```

View the rule you just created.

```
admin4@gordon-freeman:~$ sudo iptables -vL OUTPUT
Chain OUTPUT (policy ACCEPT 0 packets, 0 bytes)
  pkts bytes target     prot opt in     out      source               destination
    2   168  DROP       icmp  --  any    enp2s0   anywhere             anywhere
```

Flush all the rules and execute the ping command again.

```
admin4@gordon-freeman:~$ sudo iptables -F
admin4@gordon-freeman:~$ sudo ping -c 2 www.google.com
PING www.google.com (142.251.220.4) 56(84) bytes of data.
64 bytes from pnbomb-ay-in-f4.1e100.net (142.251.220.4): icmp_seq=1 ttl=117 time=1.99 ms
64 bytes from pnbomb-ay-in-f4.1e100.net (142.251.220.4): icmp_seq=2 ttl=117 time=2.02 ms

--- www.google.com ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1001ms
rtt min/avg/max/mdev = 1.989/2.003/2.017/0.014 ms
```

Now create another rule that will drop ICMP packets that originate from a specific unwanted IP address

```
admin4@gordon-freeman:~$ sudo iptables -A INPUT -i enp2s0 -p icmp --source 172.16.0.44 -j DROP
```

Experiment 7 :- To control network traffic using iptables rules.

Instead of flushing all the rules in your tables. Delete only the rule you created above. To do this you need to know the rule number. Using the line-number that matches the rule in you want to delete, you can delete the specific rule

```
admin4@gordon-freeman:~$ sudo iptables -vL INPUT --line-numbers
Chain INPUT (policy ACCEPT 0 packets, 0 bytes)
num  pkts bytes target     prot opt in     out      source          destination
1      0     0 DROP       icmp --  enp2s0 any    172.16.0.44      anywhere
admin4@gordon-freeman:~$ sudo iptables -D INPUT 1
```

To filter other kinds of traffic, restart ftp

```
admin4@gordon-freeman:~$ sudo service vsftpd start
admin4@gordon-freeman:~$ sudo service vsftpd restart
```

Block FTP from partner

```
admin4@gordon-freeman:~$ sudo iptables -A INPUT -i enp2s0 -s www.google.com -p tcp --dport 21 -j DROP
admin4@gordon-freeman:~$ sudo iptables -A INPUT -s www.google.com -p tcp --dport 80 -j DROP
```

Questions

**1. What option is needed to get a more verbose version of this command
iptables -L -t nat?**

Ans : iptables -L -t nat -v

2. What is the command to display the rules under the OUTPUT chain?

Answer : iptables -L OUTPUT

3. What port does the ftp service “normally” listen on?

Answer : Port 21

4. What is the command to create a chain called “mynat-chain” under the nat table?

Experiment 7 :- To control network traffic using iptables rules.

Answer : `iptables -t nat -N mynat-chain`

5. Research online and list the names of some easier to use tools or applications that can be used to manage the firewall sub-system on Linux based systems.

Answer :

- UFW (Uncomplicated Firewall) – user-friendly frontend for iptables
 - Firewalld – dynamic firewall manager with zones
 - Gufw – graphical interface for UFW
 - Shorewall (Shoreline Firewall) – high-level tool for configuring iptables
 - CSF (ConfigServer Security & Firewall) – advanced firewall UI for servers
-

Conclusion

In this experiment, I explored the functionality and practical application of iptables for managing network traffic and enhancing system security. By configuring various rules and chains, I learned how to control incoming and outgoing packets, implement firewall policies, and prevent unauthorized access. The hands-on experience helped reinforce theoretical knowledge of networking and security, and highlighted the importance of rule order and chain priority in effective firewall configuration.